

# NBA Player Performance Prediction

---

**Name:** Shefali Khera

**Roll no:** 22csu426

**Affiliation:** The NorthCap University





# INTRODUCTION

## **Background:**

- The NBA is rich with player statistics, and analyzing these stats can help predict future performance.
- Machine learning models can play a critical role in generating predictions based on past performance, helping teams optimize their strategies.

## **Problem Statement:**

- How can we use historical data from the 2023-24 season to predict future player performance metrics, like points per game, rebounds?

## **Objective:**

- Build a machine learning model to predict key performance indicators (KPIs) like points, assists, and rebounds for NBA players based on their 2023-24 season stats that help teams make informed decisions about player utilization, game strategy, and player development

# Dataset Description

## Dataset Source:

- The dataset was sourced from Kaggle and contains comprehensive NBA player statistics for the 2023-24 season, including playoff data.

## Data Overview:

- The dataset includes 68 columns and 52218 rows.
- The dataset includes categorical data like PLAYER\_NAME, TEAM\_NAME, LOCATION, and Conference, and numerical data such as FGM, MIN\_x, REB, PTS, AST, 3PA, 3PM, FTA, and FTM, which quantify player performance.
- The dataset consists of regular-season and playoff data, with no missing values or duplicate rows.

## Key Features:

- Player Stats: Minutes Played, Field Goals Made, Rebounds, Assists, etc.
- Team Stats: Team name, opponent stats such as points in the paint, and overall team rank.

# Preprocessing and Data Cleaning

## **Data Cleaning:**

- No missing values were detected, simplifying preprocessing.
- No duplicate rows were found (dups. sum() = 0)

## **Outlier Removal:**

- The Interquartile Range (IQR) calculated the bounds, capping values outside this range.

## **Feature Engineering:**

- New features like PER (Player Efficiency Rating) were created from existing statistics.
- Text vectorization on the Comments column to convert textual data into numerical features for machine learning analysis.

## **Data Transformation:**

- Categorical variables such as Conference were prepared for machine learning models using one-hot encoding.
- Log transformation to reduces skewness by converting values into their logarithmic form

# DUPLICATE VALUES

```
[ ] dups = df.duplicated()
    print('Number of duplicate rows = %d' % (dups.sum()))

    df[dups]
```

Number of duplicate rows = 0

PLAYER_NAME	PLAYER_ID	TEAM_NAME_x	LOCATION	MIN_x	FGM	FGA	FG_PCT	FG3M	FG3A	...	DEF_RATING_RANK	DREB_RANK	DREB_PCT_RANK	STL_RANK	BLK_RANK	OPP_PTS_OFF_TOV_RANK
-------------	-----------	-------------	----------	-------	-----	-----	--------	------	------	-----	-----------------	-----------	---------------	----------	----------	----------------------

0 rows x 68 columns


# NULL VALUES

```
df.isnull().sum()
```

	0
PLAYER_NAME	0
PLAYER_ID	0
TEAM_NAME_x	0
LOCATION	0
MIN_x	0
...	...
OPP_PTS_OFF_TOV_RANK	0
OPP_PTS_2ND_CHANCE_RANK	0

# OUTLIER DETECTION


```
[ ] df.boxplot(column=['FGM'])
    plt.show()
```



```
[ ] def remove_outlier(col):
    sorted(col)
    Q1,Q3=col.quantile([0.25,0.75])
    IQR=Q3-Q1
    lower_range= Q1-(1.5 * IQR)
    upper_range= Q3+(1.5 * IQR)
    return lower_range, upper_range

[ ] lrincome,urincome=remove_outlier(df['FGM'])
    df['FGM']=np.where(df['FGM']>urincome,urincome,df['FGM'])
    df['FGM']=np.where(df['FGM']<lrincome,lrincome,df['FGM'])
```

```
df.boxplot(column=['FGM'])
plt.show()
```



# ONE-HOT ENCODING

```
from sklearn.preprocessing import OneHotEncoder
encoder = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
X_encoded = encoder.fit_transform(df[['Conference']])

print("Shape of X_encoded:", X_encoded.shape)
print("First few rows of X_encoded:\n", X_encoded[:5])
```

```
→ Shape of X_encoded: (2418, 2)
First few rows of X_encoded:
[[1. 0.]
 [1. 0.]
 [1. 0.]
 [1. 0.]
 [0. 1.]]
```

# LOG TRANSFORMATION

```
import numpy as np
df['log_PTS'] = np.log(df['PTS'] + 1)

print(df[['PTS', 'log_PTS']])
```

```
→
```

	PTS	log_PTS
0	4.11	1.631199
1	4.86	1.768150
2	2.75	1.321756
3	2.17	1.153732
4	3.63	1.532557
...	...	...
2413	29.83	3.428488
2414	7.57	2.148268
2415	1.17	0.774727
2416	1.50	0.916291
2417	1.67	0.982078

```
[2418 rows x 2 columns]
```

- **OneHotEncoder** transforms categorical data (like 'Conference') into numerical format using binary vectors
- **Log transformation** `np.log(df['PTS'] + 1)` calculates the natural logarithm of the 'PTS' column, adding 1 to avoid log(0).

# Exploratory Data Analysis (EDA)

## Visualizations


- A correlation matrix was generated using numeric columns.
- A heatmap visually represented the strength of these correlations.
- A count plot showed the distribution of players across different conferences.
- Radar charts to compare player statistics across different dimensions
- Create word clouds to visualize the most common words used like player comments.

## Key Insights

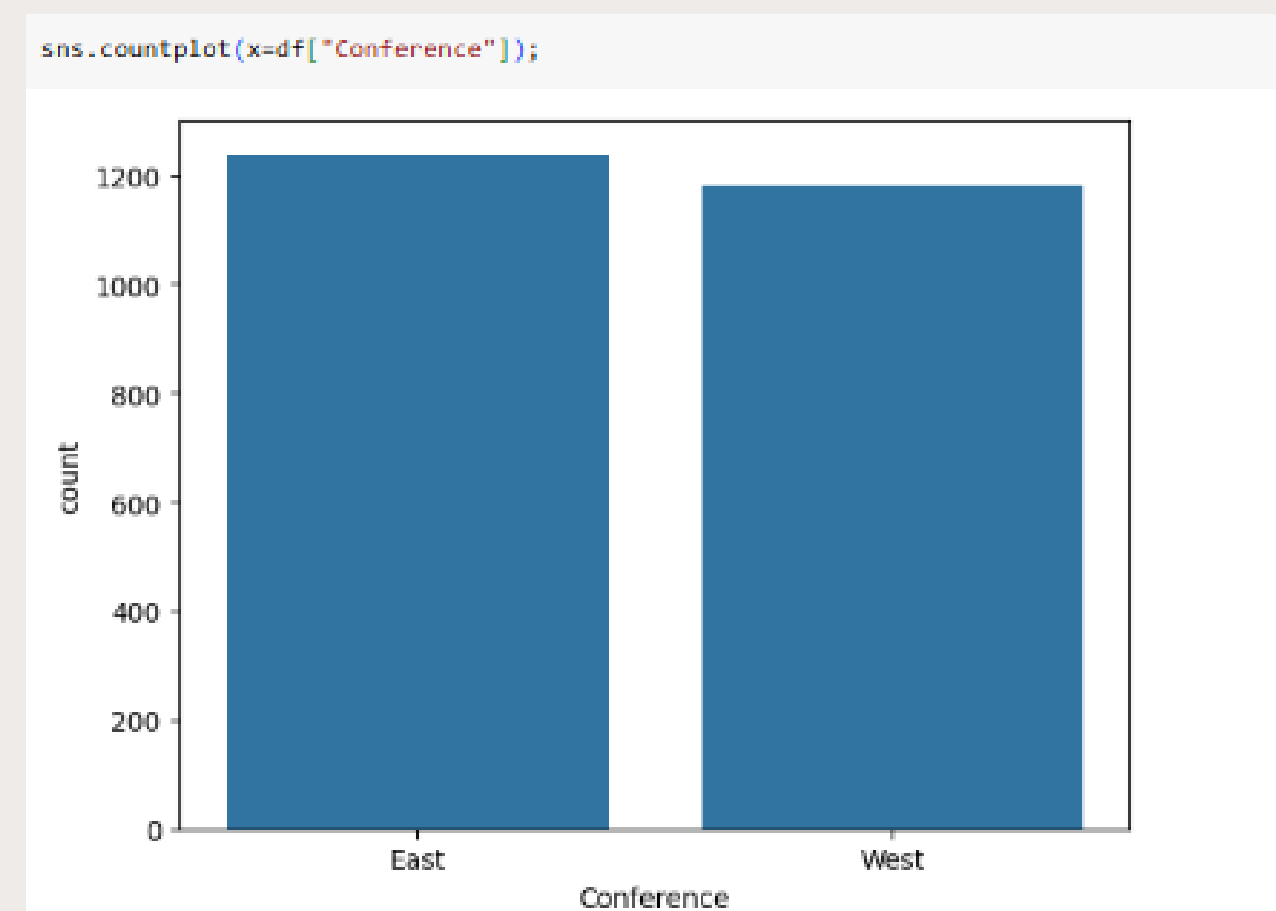
- The correlation heatmap revealed a strong positive correlation between assists (AST) and points (PTS). Players who assist more tend to score more points.
- A significant correlation was found between rebounds (REB) and minutes played (MIN\_x), indicating that players on the court longer secure more rebounds.
- These insights can guide strategic decisions in team management, player development, and understanding player dynamics.



# COUNT PLOT




# HIST PLOT



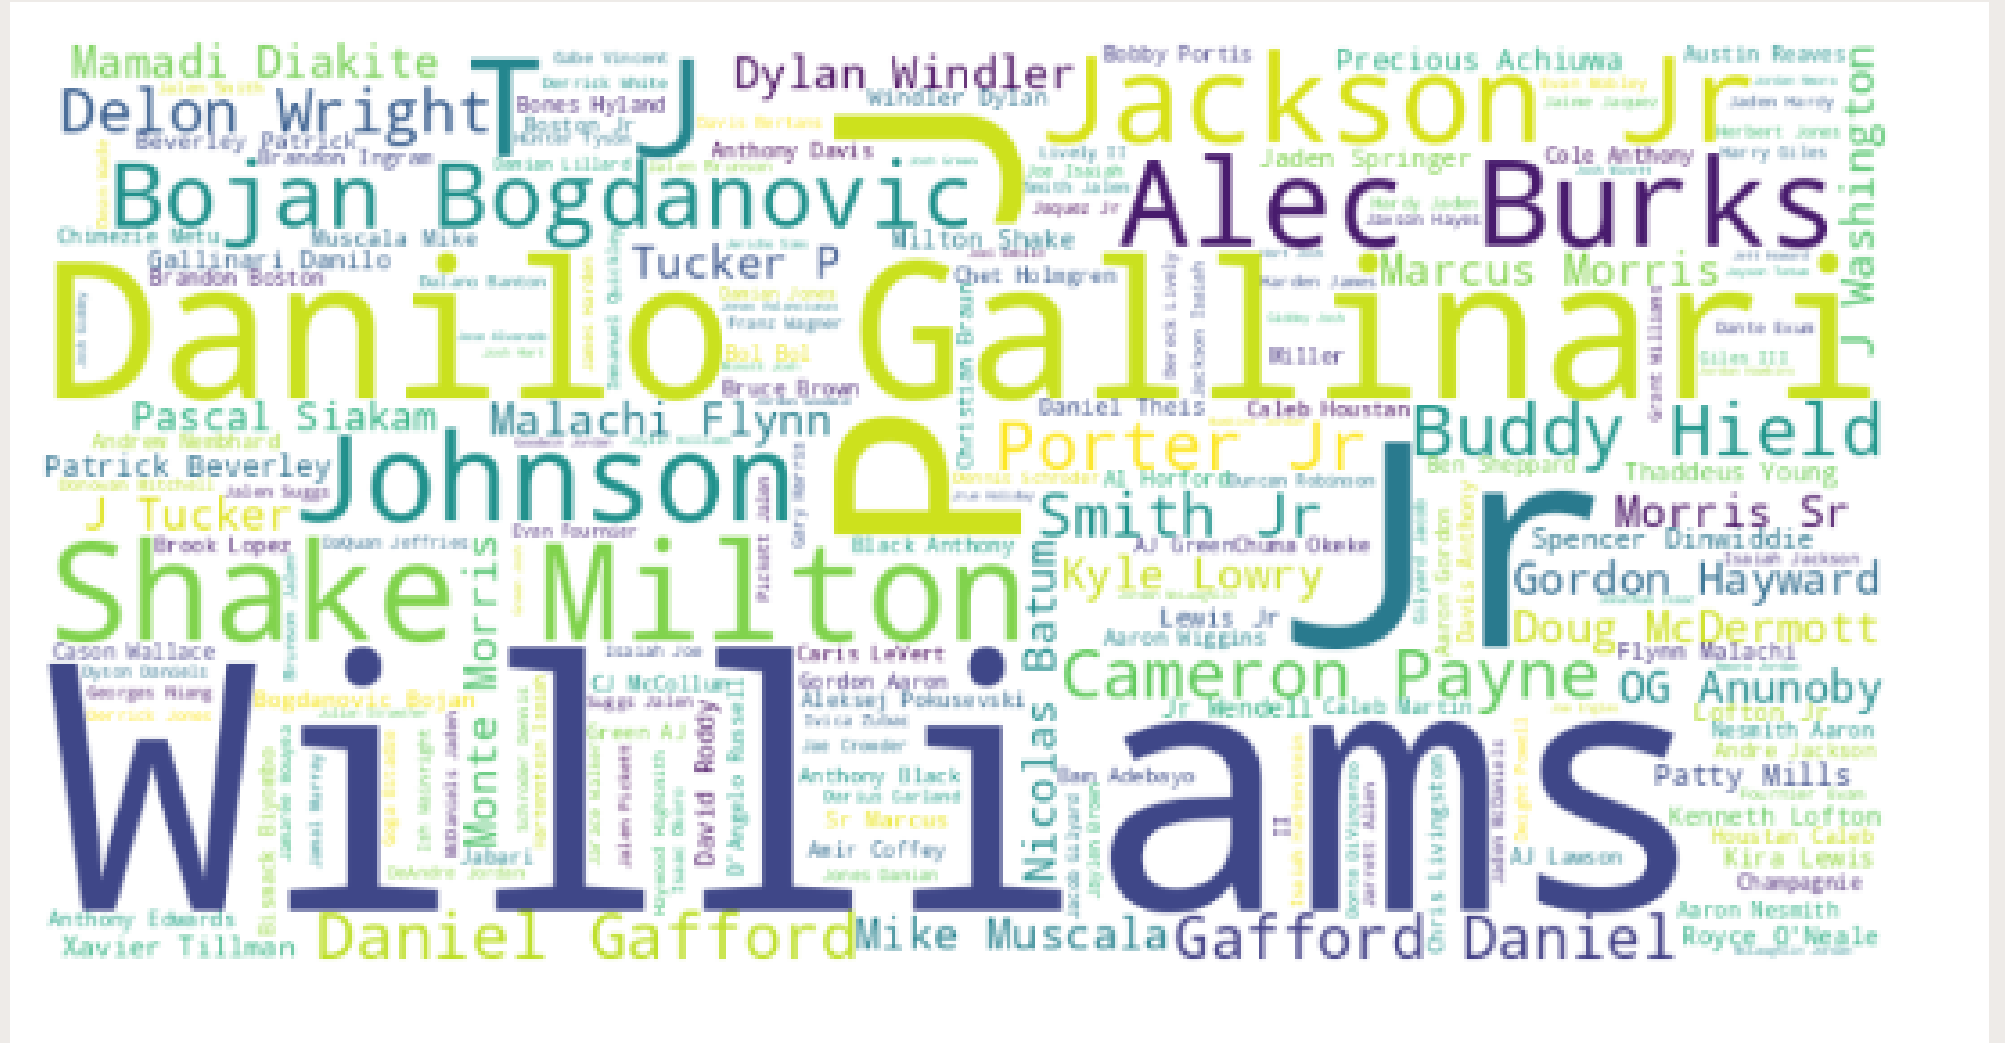

- The **correlation matrix** shows the relationships between numerical variables, while the **count plot** displays the frequency of each conference and the **histogram** shows the distribution of field goal made.

# CORRELATION MATRIX






# RADAR CHART




- The **radar chart** visualizes the average statistics (PTS, AST, REB, STL, BLK) of several players. Each player's stats are represented by a polygon, allowing for a comparison of their performance across various metrics. It helps in understanding the strengths and weaknesses of individual players relative to others.

# SCATTER PLOT

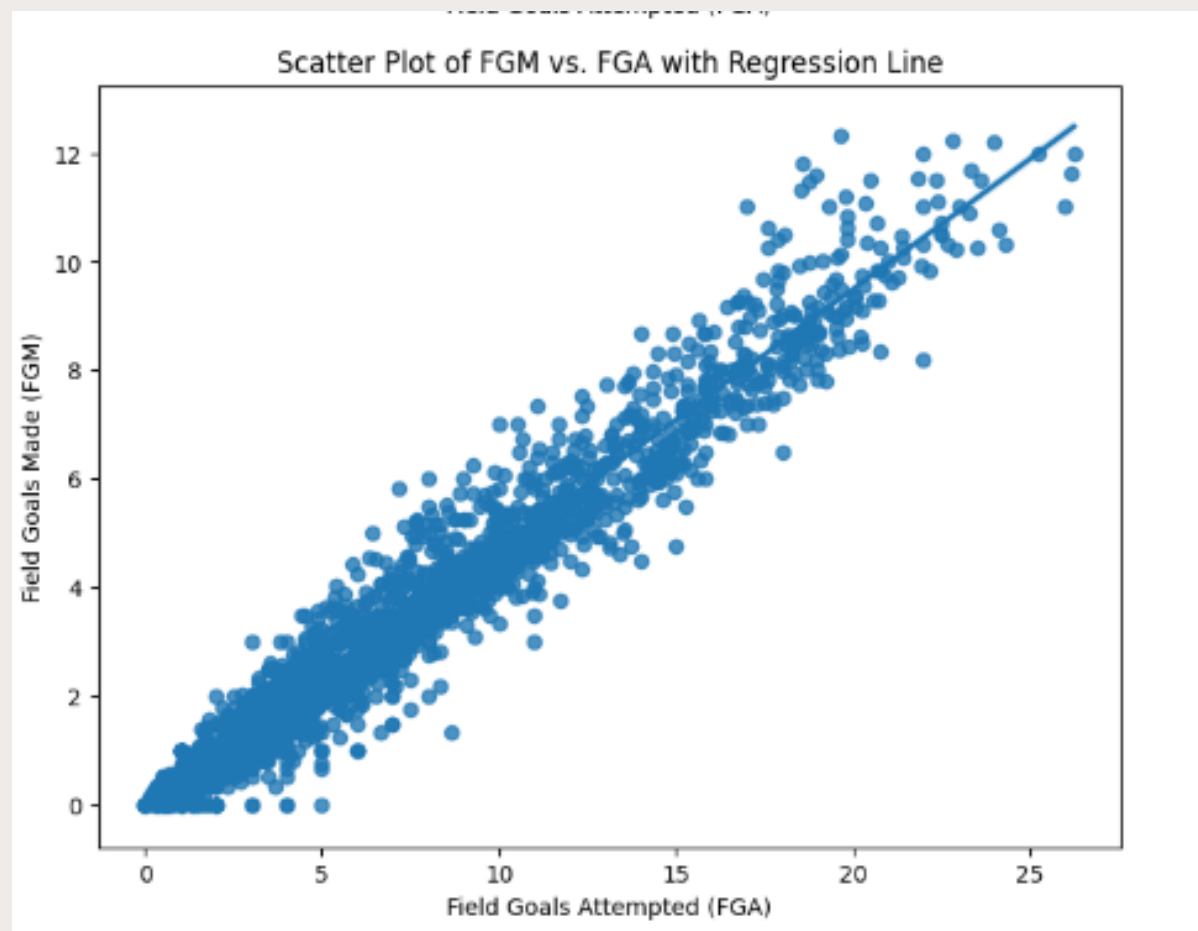


- **Scatter Plot:** Shows the relationship between two continuous variables (e.g., FGA and FGM).
- **Box Plot:** Displays the distribution of a continuous variable (e.g., PTS) across different categories (e.g., conferences). It provides information on the median, quartiles, and outliers.
- **Violin Plot:** Similar to box plot, but also shows the probability density of the data at different values, providing a more complete view of the distribution.

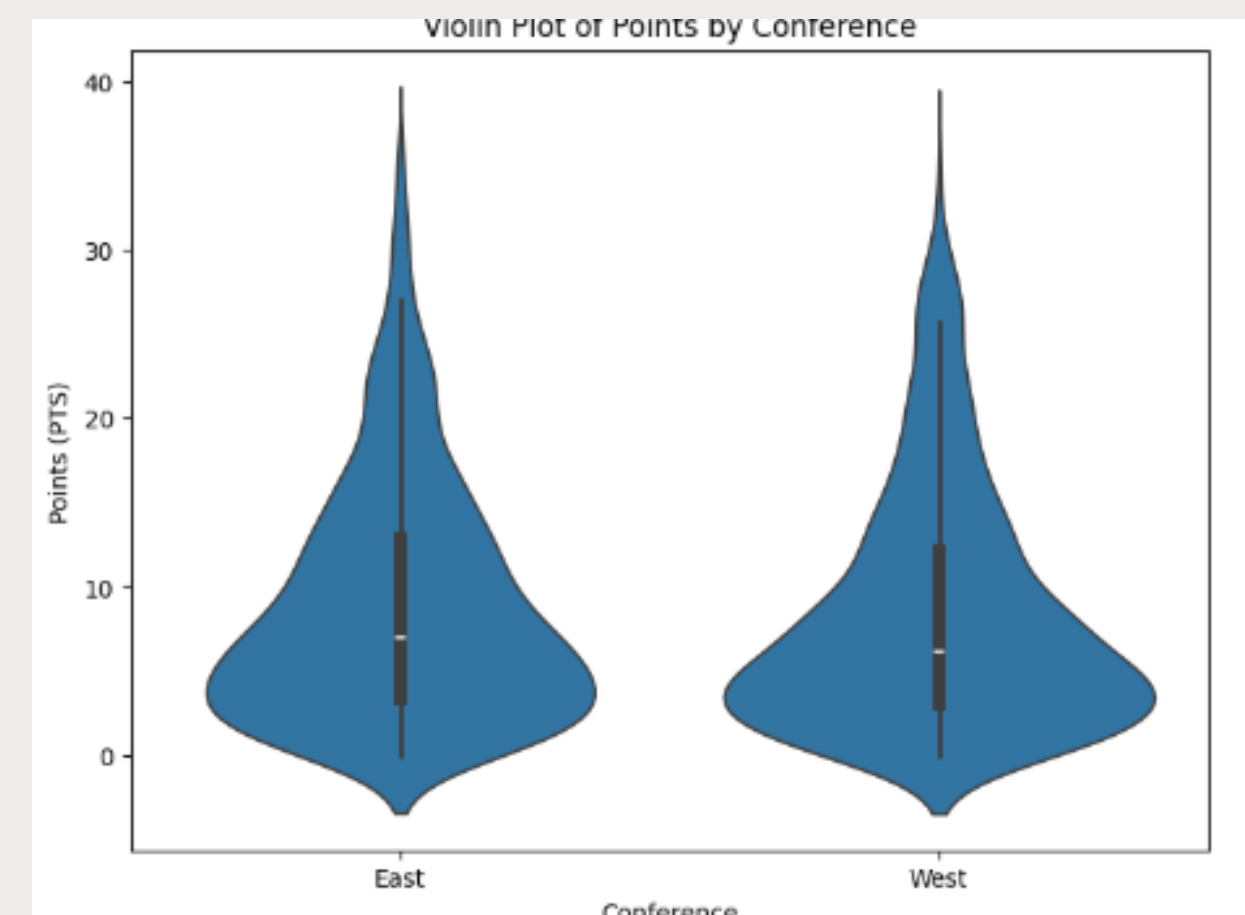
# SCATTER PLOT



# SCATTER PLOT



# VIOLIN PLOT



# Machine Learning Models

## Regression Models

- Linear Regression:

Application: Predict continuous outcomes such as points scored (PTS) based on features like assists (AST), rebounds (REB), and field goals made (FGM).

- XGBoost (for Regression):

Application: Predict continuous performance metrics like Player Efficiency Rating (PER) using player statistics.

## Classification Models

- Random Forest:

Application: Classify players based on performance categories, such as whether a player will score above a certain threshold in a game.

- XGBoost (for Classification):

Application: Predict categorical outcomes, such as whether a player will be an All-Star based on their performance metrics.

- Neural Networks :

Application: Classify player performance into categories like "high scorer" or "role player" based on their statistics.

# PREDICTION FOR CLASSIFICATION

## PREDICTION

```
[ ] sample_data = [[30, 8, 2, 5, 3, 6]]

[ ] sample_df = pd.DataFrame(sample_data, columns=['MIN_x', 'FGM', 'FG3M', 'FTM', 'OREB', 'DREB_x'])

[ ] predicted_conference = nb_model.predict(sample_df)
   print("Predicted Conference:", predicted_conference[0])
```

➡ Predicted Conference: East

- This code uses a classification model with features (`MIN\_x`, `FGN`, `FG3M`, `FTM`, `OREB`, `DREB\_x`) to predict the **conference** (East or West) and prints the predicted result.

# PREDICTION FOR REGRESSION

## PREDICTION

```
[76] sample_data = pd.DataFrame({
      'MIN_x': [200],
      'FGM': [0.5],
      'FG3M': [2],
      'FTM': [5],
      'OREB': [8],
      'DREB_x': [20]
    })
   predicted_pts = regressor.predict(sample_data)

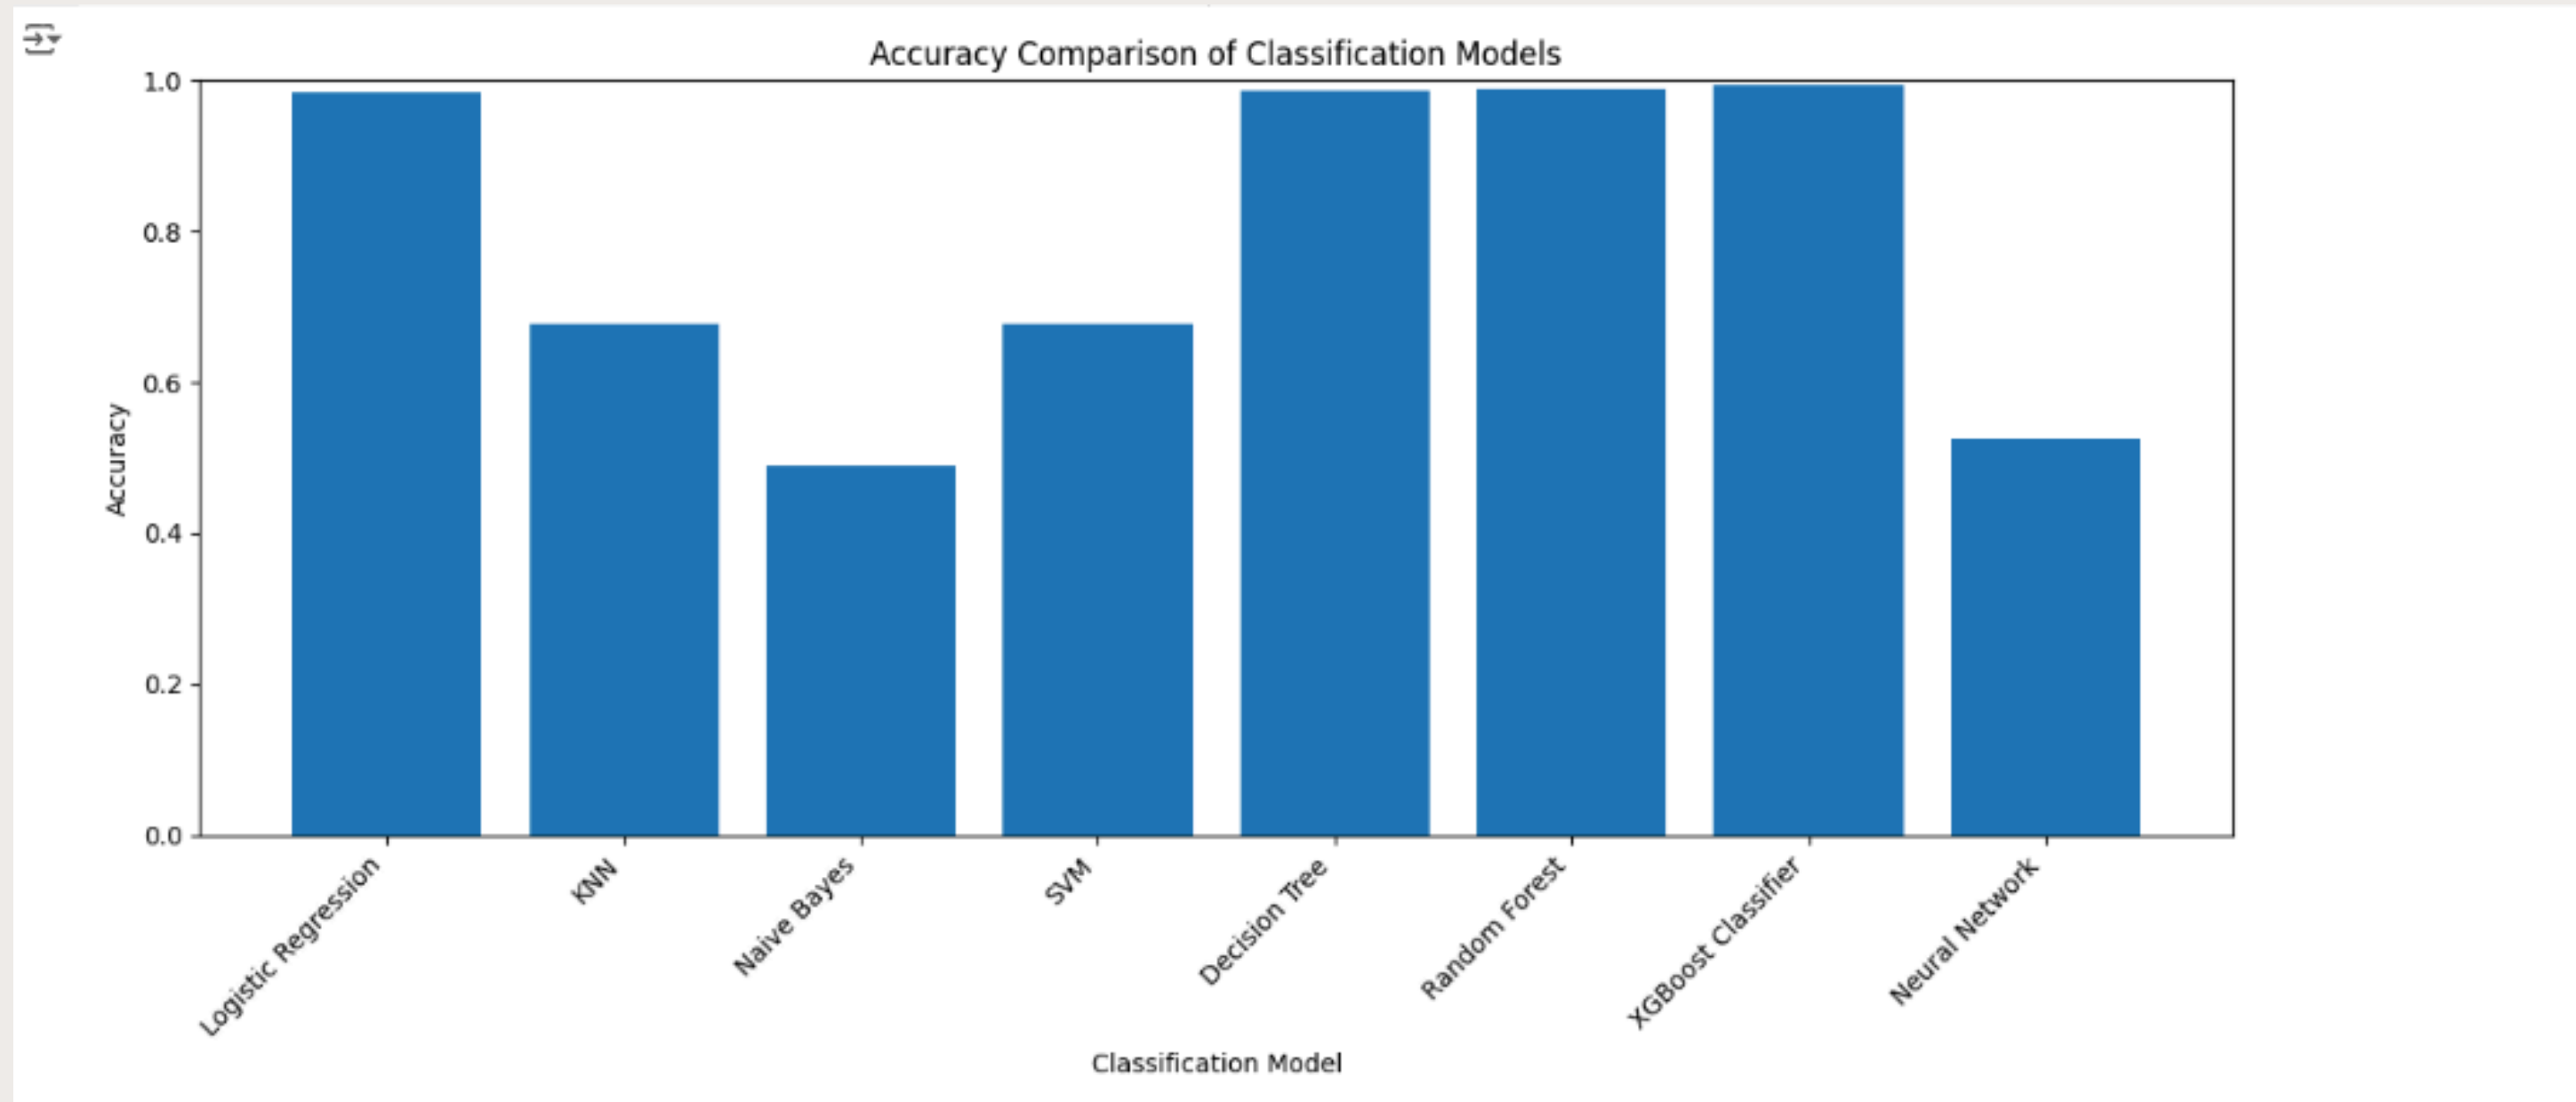
   print("Predicted PTS:", predicted_pts[0])
```

➡ Predicted PTS: 8.000669102757023

- This code creates sample data with features (`MIN\_x`, `FGN`, `FG3M`, `FTM`, `OREB`, `DREB\_x`) and uses a trained model (`regressor`) to predict **points** (`PTS`), then prints the predicted result.

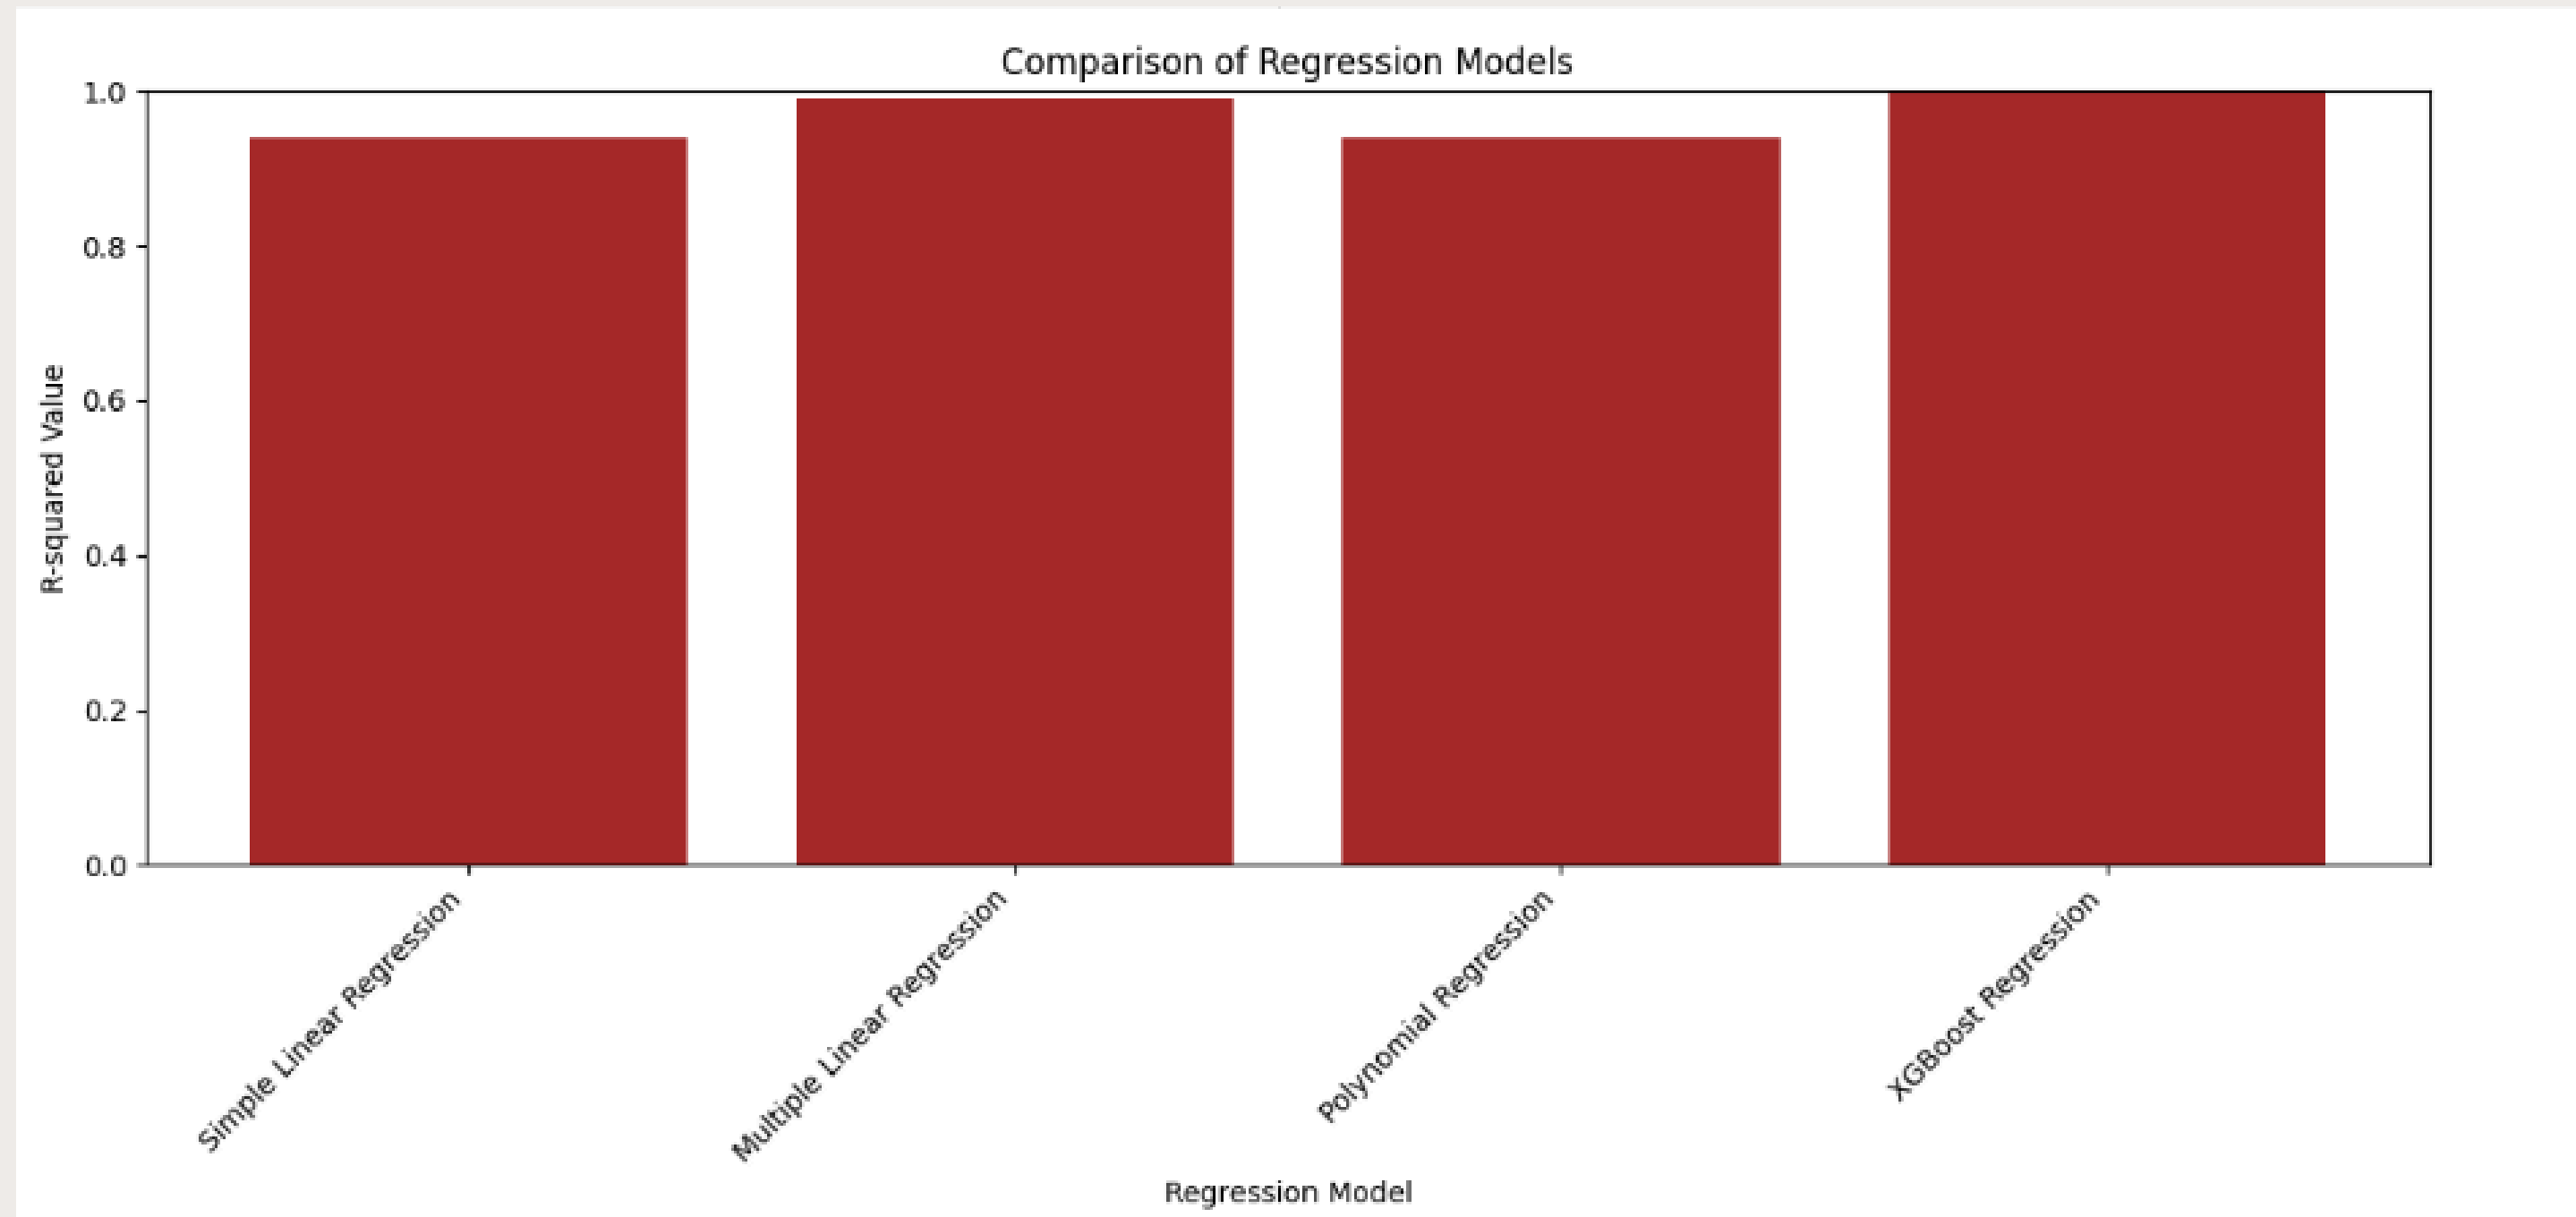


# Classification Model Analysis



- **XGBoost Classifier** and **Random Forest** have the highest accuracy, indicating that they are the best-performing models for predicting the 'Conference' variable. Neural Network and Naive Bayes show lower accuracy, suggesting they may not be the most suitable choices for this task.

# Regression Model Analysis



- **XGBoost Regression** demonstrates the highest R-squared value. This implies that it explains the variance in the 'PTS' variable (the target for regression) the best. Multiple Linear Regression also shows a high R-squared, meaning that it fits the data well.

# Feature Selection

Feature selection using `SelectKBest` with `f_classif` (ANOVA F-value) offers several benefits:

1. **Improved Model Performance:** Selecting the most relevant features reduces noise and irrelevant information that can negatively impact model performance. This often leads to higher accuracy, precision, recall, and other evaluation metrics.
2. **Reduced Overfitting:** Using fewer features can help prevent overfitting, which occurs when a model performs well on training data but poorly on unseen data. By focusing on the most informative features, we create a simpler and more generalizable model.
3. **Faster Training:** With fewer features, the model training process is faster and requires fewer computational resources.
4. **Enhanced Interpretability:** A model with a smaller set of features can be easier to interpret and understand. You can gain insights into the underlying relationships between the selected features and the target variable.

- We are using feature selection to identify the **top 10** features that are most relevant to predicting the `Conference` variable.
- The **`f\_classif`** function is used for feature selection in classification problems.
- **`SelectKBest`** selects the top k features based on their statistical significance.

```
⇒ XGBoost Classifier Accuracy (Without Feature Selection): 0.993801652892562
   /usr/local/lib/python3.10/dist-packages/sklearn/feature_selection/_univariate_selection.py:112: UserWarning: Features [23]
     warnings.warn("Features %s are constant." % constant_features_idx, UserWarning)
   /usr/local/lib/python3.10/dist-packages/sklearn/feature_selection/_univariate_selection.py:113: RuntimeWarning: invalid va
     f = msb / msw
   XGBoost Classifier Accuracy (With Feature Selection): 1.0
   Difference in Accuracy: 0.006198347107438051
```

- The XGBoost classifier in the code achieves an accuracy of approximately 0.9938. This signifies a very high level of accuracy, indicating that the model is performing exceptionally well in predicting the 'Conference' variable. The model's performance is further enhanced by feature selection, resulting in potentially a slightly higher accuracy.



**THANK YOU**